



# TensorFlow TUTORIAL


Le Anh Tu



Faculty of Electrical & Electronics Engineering,  
Department of Electronics & Telecommunications Engineering,  
Ho Chi Minh City University of Technology, Vietnam

July, 2017

# Outline



**Introduction**

**Basic Usage**

**Build Model**

**Other Utilities**

# WHAT IS TENSORFLOW ?



- Tensorflow is a deep learning library recently open-sourced by Google.
- But what does it actually do?
  - Tensorflow provides primitives for defining functions on tensors and automatically computing their derivatives

# Outline



**Introduction**

**Basic Usage**

**Build Model**

**Other Utilities**

# BASIC USAGE

example 1:

- `a = tf.constant(2)`
- `b = tf.constant(3)`
- `c = a + b`

example 2:

- `x = tf.placeholder(tf.int16)`
- `y = tf.placeholder(tf.int16)`
- `z = tf.add(x,y)`

example 3:

```
input = tf.placeholder(tf.float32)
```

```
weight = tf.constant([[2.],[2.]])
```

```
bias = tf.constant(3.0)
```

```
output = tf.add(tf.matmul(input,weight),  
bias)
```

# LOAD DATA & VARIABLES INITIALIZER

- `from tensorflow.examples.tutorials.mnist import input_data`
- `mnist = input_data.read_data_sets("mnist_data/", one_hot = True)`
- `learning_rate = 0.001`
- `training_epochs = 15`
- `batch_size = 100`
- `display_step = 1`
- `n_input = 784 #shape:28x28`
- `n_classes = 10 #0-9 digits`
- `x = tf.placeholder('float', [None, n_input])`
- `y = tf.placeholder('float', [None, n_classes])`

# Outline



**Introduction**

**Basic Usage**

**Build Model**

**Other Utilities**

# NEURAL NETWORK

- `n_hidden_1 = 256`
- `n_hidden_2 = 256`
- `def neural_network(x, weights, biases):`
  - `layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])`
  - `layer_1 = tf.nn.relu(layer_1)`
  - `layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])`
  - `layer_2 = tf.nn.relu(layer_2)`
  - `out_layer = tf.matmul(layer_2, weights['out']) + biases['out']`
  - `return out_layer`
- `weights = {`
  - `'h1': tf.variable(tf.random_normal([n_input, n_hidden_1])),`
  - `'h2': tf.variable(tf.random_normal([n_hidden_1, n_hidden_2])),`
  - `'out': tf.variable(tf.random_normal([n_hidden_2, n_classes]))`
- `}`
- `biases = {`
  - `'b1': tf.variable(tf.random_normal([n_hidden_1])),`
  - `'b2': tf.variable(tf.random_normal([n_hidden_2])),`
  - `'out': tf.variable(tf.random_normal([n_classes]))`
- `}`
- `y_pred = neural_network(x, weights, biases)`
- `cost =`  
`tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(logits=y_pred, labels=y))`
- `optimizer =`  
`tf.train.adam_optimizer(learning_rate=learning_rate).minimize(cost)`



# NEURAL NETWORK

- with `tf.session()` as `sess`:
  - `sess.run(tf.global_variables_initializer())`
  - `#training loop`
  - for `epoch` in `range(training_epochs)`:
    - `avg_cost = 0.`
    - `total_batch = int(mnist.train.num_examples / batch_size)`
    - `#loop over all batches:`
    - for `i` in `range(total_batch)`:
      - `batch_x, batch_y = mnist.train.next_batch(batch_size)`
      - `_, c = sess.run([optimizer, cost], feed_dict={x: batch_x, y: batch_y})`
      - `#comput average loss`
      - `avg_cost += c / total_batch`
    - if `epoch % display_step == 0`:
      - `print('epoch:', '%04d' % (epoch + 1), 'cost =', \`  
`{:.9f}'.format(avg_cost))`
  - `print('optimization finished')`
  - `#test mode`
  - `correct_prediction = tf.equal(tf.argmax(y_pred, 1), tf.argmax(y, 1))`
  - `accuracy = tf.reduce_mean(tf.cast(correct_prediction, 'float'))`
  - `print('accuracy:', accuracy.eval({x: mnist.test.images, y: mnist.test.labels}))`

```
Epoch: 0001 cost= 183.146648452
Epoch: 0002 cost= 43.430120397
Epoch: 0003 cost= 26.733731529
Epoch: 0004 cost= 18.591381240
Epoch: 0005 cost= 13.488394500
Epoch: 0006 cost= 9.907654405
Epoch: 0007 cost= 7.318386168
Epoch: 0008 cost= 5.641316663
Epoch: 0009 cost= 4.140372237
Epoch: 0010 cost= 3.029535837
Epoch: 0011 cost= 2.293382505
Epoch: 0012 cost= 1.781092959
Epoch: 0013 cost= 1.289214330
Epoch: 0014 cost= 1.023641920
Epoch: 0015 cost= 0.802592557
Optimization Finished!
Accuracy: 0.949
```

# CONVOLUTIONAL NEURAL NETWORK

- `training_inters = 200000`
- `dropout = 0.75`
- `def conv2d(x,w,b,strides=1):`
  - `x = tf.nn.conv2d(x,w, strides = [1,strides,strides,1],padding = 'same')`
  - `x = tf.nn.bias_add(x,b)`
  - `return tf.nn.relu(x)`
- `def maxpool2d(x, k=2):`
  - `return tf.nn.max_pool(x, ksize = [1,k,k,1], strides = [1,k,k,1], padding = 'same')`
- `def conv_net(x,weights, biases, dropout):`
  - `x = tf.reshape(x, shape = [-1,28,28,1])`
  - `conv1 = conv2d(x, weights['wc1'], biases['bc1'])`
  - `conv1 = maxpool2d(conv1, k=2)`
  - `conv2 = conv2d(conv1 , weights['wc2'], biases['bc2'])`
  - `conv2 = maxpool2d(conv2, k=2)`
  - `fc1 = tf.reshape(conv2, [-1, weights['wd1'].get_shape().as_list()[0]])`
  - `fc1 = tf.add(tf.matmul(fc1, weights['wfl']), biases['bfl'])`
  - `fc1 = tf.nn.relu(fc1)`
  - `fc1 = tf.nn.dropout(fc1, dropout)`
  - `out = tf.add(tf.matmul(fc1, weights['out']), biases['out'])`
  - `return out`

# CONVOLUTIONAL NEURAL NETWORK

- weights = {
  - 'wc1': tf.Variable(tf.random\_normal([5,5,1,32])),
  - 'wc2': tf.Variable(tf.random\_normal([5,5,32,64])),
  - 'wf1': tf.Variable(tf.random\_normal([7\*7\*64, 1024])),
  - 'out': tf.Variable(tf.random\_normal([1024, n\_classes]))
- }
- biases = {
  - 'bc1': tf.Variable(tf.random\_normal([32])),
  - 'bc2': tf.Variable(tf.random\_normal([64])),
  - 'bf1': tf.Variable(tf.random\_normal([1024])),
  - 'out': tf.Variable(tf.random\_normal([n\_classes]))
- }

```
Iter 167680, Minibatch Loss= 51.544586, Training Accuracy= 0.97656
Iter 168960, Minibatch Loss= 16.598854, Training Accuracy= 0.99219
Iter 170240, Minibatch Loss= 102.296074, Training Accuracy= 0.97656
Iter 171520, Minibatch Loss= 126.573250, Training Accuracy= 0.97656
Iter 172800, Minibatch Loss= 50.975868, Training Accuracy= 0.98438
Iter 174080, Minibatch Loss= 13.026268, Training Accuracy= 0.99219
Iter 175360, Minibatch Loss= 61.759331, Training Accuracy= 0.98438
Iter 176640, Minibatch Loss= 104.688667, Training Accuracy= 0.97656
Iter 177920, Minibatch Loss= 117.068787, Training Accuracy= 0.96875
Iter 179200, Minibatch Loss= 68.032127, Training Accuracy= 0.97656
Iter 180480, Minibatch Loss= 55.089828, Training Accuracy= 0.99219
Iter 181760, Minibatch Loss= 12.846817, Training Accuracy= 0.98438
Iter 183040, Minibatch Loss= 19.284180, Training Accuracy= 0.98438
Iter 184320, Minibatch Loss= 415.337708, Training Accuracy= 0.95312
Iter 185600, Minibatch Loss= 0.000000, Training Accuracy= 1.00000
Iter 186880, Minibatch Loss= 133.618011, Training Accuracy= 0.96875
Iter 188160, Minibatch Loss= 314.425293, Training Accuracy= 0.96875
Iter 189440, Minibatch Loss= 97.414436, Training Accuracy= 0.97656
Iter 190720, Minibatch Loss= 22.250843, Training Accuracy= 0.98438
Iter 192000, Minibatch Loss= 24.891357, Training Accuracy= 0.99219
Iter 193280, Minibatch Loss= 223.365280, Training Accuracy= 0.96094
Iter 194560, Minibatch Loss= 85.471329, Training Accuracy= 0.98438
Iter 195840, Minibatch Loss= 35.877510, Training Accuracy= 0.99219
Iter 197120, Minibatch Loss= 127.821899, Training Accuracy= 0.97656
Iter 198400, Minibatch Loss= 20.495216, Training Accuracy= 0.99219
Iter 199680, Minibatch Loss= 194.098282, Training Accuracy= 0.97656
Optimization Finished!
Testing Accuracy: 0.976563
```

# RECURRENT NEURAL NETWORK

- `from tensorflow.contrib import rnn`
- `n_input = 28 #img shape 28x28`
- `n_steps = 28 #timesteps`
- `n_hidden = 128 #hidden layer num of features`
- `n_classes = 10 #0-9 digits`
- `x = tf.placeholder('float', [none, n_steps, n_input])`
- `y = tf.placeholder('float', [none, n_classes])`
- `weights = {`
  - `'out': tf.variable(tf.random_normal([n_hidden, n_classes]))`
- `}`
- `biases = {`
  - `'out': tf.variable(tf.random_normal([n_classes]))`
- `}`
- `def rnn(x, weights, biases):`
  - `x = tf.unstack(x, n_steps, 1)`
  - `lstm_cell = rnn.basicalstmcell(n_hidden, forget_bias = 1.0)`
  - `outputs, states = rnn.static_rnn(lstm_cell, x, dtype=tf.float32)`
  - `return tf.matmul(outputs[-1], weights['out']) + biases['out']`
- `y_pred = rnn(x, weights, biases)`

```
Iter 80000, Minibatch Loss= 0.121207, Training Accuracy= 0.95075
Iter 81920, Minibatch Loss= 0.064821, Training Accuracy= 0.97656
Iter 83200, Minibatch Loss= 0.114697, Training Accuracy= 0.97656
Iter 84480, Minibatch Loss= 0.115060, Training Accuracy= 0.96875
Iter 85760, Minibatch Loss= 0.162514, Training Accuracy= 0.95312
Iter 87040, Minibatch Loss= 0.173857, Training Accuracy= 0.95312
Iter 88320, Minibatch Loss= 0.107287, Training Accuracy= 0.96094
Iter 89600, Minibatch Loss= 0.121720, Training Accuracy= 0.95312
Iter 90880, Minibatch Loss= 0.112574, Training Accuracy= 0.97656
Iter 92160, Minibatch Loss= 0.138351, Training Accuracy= 0.96094
Iter 93440, Minibatch Loss= 0.209882, Training Accuracy= 0.93750
Iter 94720, Minibatch Loss= 0.088629, Training Accuracy= 0.98438
Iter 96000, Minibatch Loss= 0.089171, Training Accuracy= 0.96094
Iter 97280, Minibatch Loss= 0.161869, Training Accuracy= 0.94531
Iter 98560, Minibatch Loss= 0.026504, Training Accuracy= 1.00000
Iter 99840, Minibatch Loss= 0.047803, Training Accuracy= 0.98438
Optimization Finished!
Testing Accuracy: 0.953125
```

# Outline



**Introduction**

**Basic Usage**

**Build Model**

**Other Utilities**

# SAVE AND RESTORE MODEL

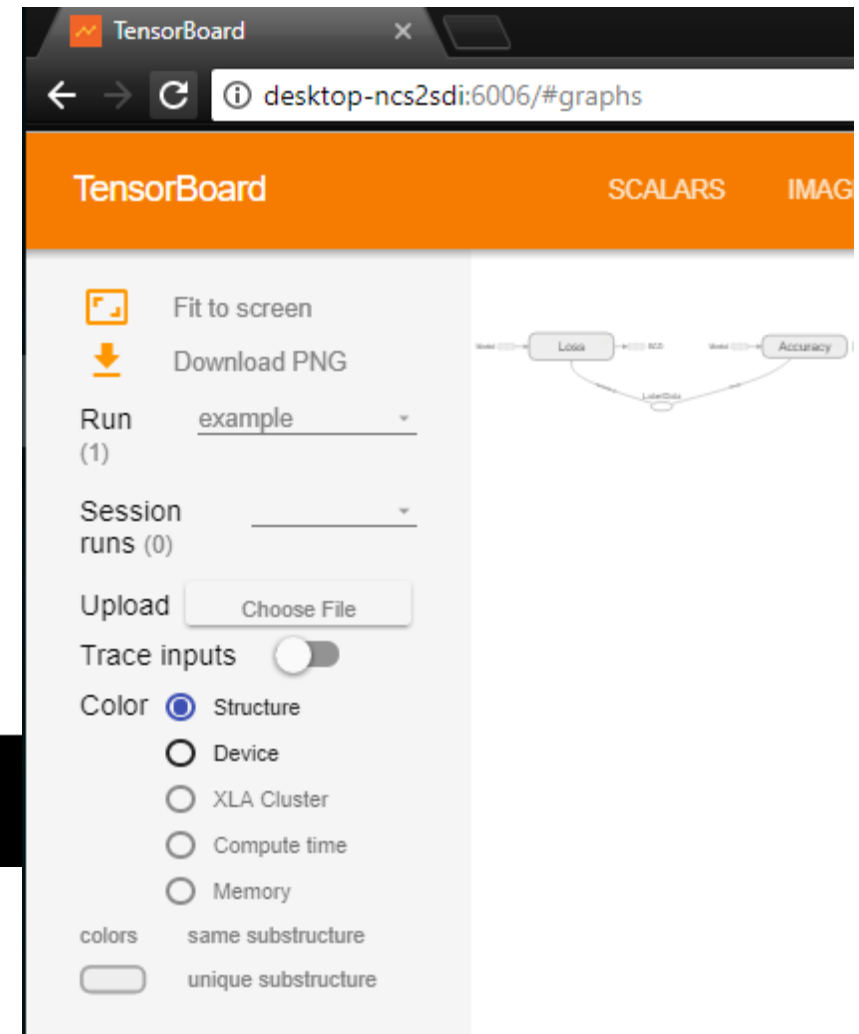
- `model_path = './model/model.ckpt'`
- `#Build your model below`
- `...`
- `#Put here before running session`
- `saver = tf.train.Saver()`
- `#If you want to save your trained model`
- `save_path = saver.save(sess,model_path)`
- `#If you want to load your saved model to reuse`
- `load_path = saver.restore(sess,model_path)`

# TENSORBOARD

- See more at
- [https://www.tensorflow.org/get\\_started/summaries\\_and\\_tensorboard](https://www.tensorflow.org/get_started/summaries_and_tensorboard)

```
tensorboard --logdir=/tmp/tensorflow_logs
```

```
(C:\Users\engleanhtu\Anaconda3) C:\Users\engleanhtu>tensorboard --logdir=/tmp/tensorflow_logs  
Starting TensorBoard b'54' at http://DESKTOP-NCS2SDI:6006  
(Press CTRL+C to quit)
```



# OPERATION ON MULTI-GPU

- # "/cpu:0":The CPU of your machine.
- # "/gpu:0":The first GPU of your machine
- # "/gpu:1":The second GPU of your machine
- #Processing Units Logs
- Log\_device\_placement = True
- #Using single GPU to compute
- With tf.device('/gpu:0'):
  - ...
- #Using CPU to compute
- With tf.device('/cpu:0'):
  - ...
- #Using multi GPU to compute
- With tf.device('/gpu:0'):
  - ...
- With tf.device('/gpu:1'):
  - ...
- With tf.device('/cpu:0'):
  - ...





# THANK YOU

**E-mail:** [eng.leanhtu@gmail.com](mailto:eng.leanhtu@gmail.com)

Code: [github.com/engleanhtu/tensorflowtutorial](https://github.com/engleanhtu/tensorflowtutorial)

*Suggestions Questions*